

# Advanced Programming Tutorial

## Worksheet 4: Resource Management

**Keywords:** raw pointers, `new`, `delete`, address sanitizer, `unique_ptr`, `std::map`, `struct`, `->`

### Prepare!

Before coming to the tutorial, and so that the time is enough to solve the exercise, revise the slides of the lecture unit “Resource management” and answer the following questions:

- ☐ What is the difference between references and pointers?
- ☐ What do the keywords `new` and `delete` do? What is a memory leak?
- ☐ What is RAI? How do smart pointers implement RAI?
- ☐ What are some ways to define arrays of fixed and of dynamic size?
- ☐ What is an iterator and where is it used?
- ☐ What is an lvalue and what is an rvalue? What does `&&` mean?
- ☐ Check that your compiler provides address sanitizers. On Linux with GCC, `g++ -fsanitize=address` should not complain about an unrecognized option.
- ☐ Remember what you did in Worksheet 2, exercise 1 (`vis_app`). Open the skeleton file `vis_app2.cpp`: Which parts of the code can you ignore, and where will you be adding your changes?
- ☐ Read through the “some new concepts” section at the end.

### Exercise 1: Matrix-vector product - A Space Odyssey

Sometimes you can’t escape your legacy. There will be times that you have to deal with lower-level code. Let’s explore how we could have implemented the same matrix-vector product from Worksheet 3 in 2001.

In older C++, we used to represent 2D matrices as pointers to 1D arrays, like this:

```
1 double * matrix = new double[matrix_rows * matrix_columns];
2 matrix[i*matrix_columns+j]
```

Doesn’t this code smell like the old wooden desk and red leather chair it was written at?

The file `matvec.cpp` contains the solution from Worksheet 3 and an alternative, legacy-style implementation `mat_times_vec()` of the matrix-vector product we implemented in `operator*`.

1. Read through the code: what differences do you observe? Which additional keywords, types, and variables do you see?
2. Important to remember is that we always need to `delete` everything that we allocate with `new`. However, it is easy to forget something. Compile your code with GCC and use the options `-fsanitize=address` and `-g` (to enable debug symbols):  
`g++ matvec.cpp -fsanitize=address -g -o matvec`  
What additional information do you get when you run the program?

3. There is a memory leak in the program and the sanitizer tell you where. Fix it!
4. **Idea:** Can you think of a better way to organize your code so that you can more easily avoid the potential memory leaks? Think of a principle we discussed in the lecture.

Note: You are of course using Git to keep track of your exercises, right? ;-) Since we will not go too far in this direction, you may want to make a separate git branch for this.

## Exercise 2: The pandemic (tracker) gets real (data)

In Worksheet 2, you already created your first pandemic tracker. For the purpose of quick development, you used pseudo-randomly generated data. Now you want to also process real data from a CSV file<sup>1</sup> However, you want to keep the dummy data to be able to easily test your implementation, without changing the code.

In this exercise, we will use a `unique_ptr` to select between the dummy and the real dataset at runtime. We will also learn how to pass the `unique_ptr` to a function and we will work with `std::map` to create views of the dataset.

Open the skeleton file `caseTrack.cpp` and:

1. Take 5-10min to read through the code and the section after this exercise. There are code parts you don't need to read right now (see guiding comments).
2. Fill the `TODO` parts in `main()`:
  - (a) Make the unique pointer `data_frame` so that it points to dummy or real data, depending on the `if` branch.
  - (b) Print the cases for each country, independent of the kind of data (dummy or real).
3. Implement `normalize_per_capita(std::unique_ptr<DataFrame>& data_frame)`. This should return a map named `cases_normalized` with the cases per 100,000 people for each country and day.
4. Call the function `normalize_per_capita` on the given `data_frame` and print the results.

---

<sup>1</sup>European Centre for Disease Prevention and Control, file formatted for our needs.

## Some new concepts

### Structures: struct

You may already know classes: that's good! You can think of structures as simplified classes (with public members by default). You don't already know what a class is? Then you can think of a struct as a collection of data, such as:

```
1 struct Color {  
2     int red;  
3     int green;  
4     int blue;  
5 };
```

We can then create a specific color and access the underlying elements as:

```
1 Color TUMBlue = {0, 101, 189};  
2 std::cout << TUMBlue.red << "\n";
```

### Arrow operator: ->

If we access our color through a pointer (raw or smart), we first need to dereference it to access the members of it:

```
1 Color TUMBlue = {0, 101, 189}; // ok, normally no temporary object needed  
2 unique_ptr<Color> myFavoriteColor = make_unique<Color>(TUMBlue);  
3 std::cout << (*myFavoriteColor).red << "\n";
```

To make our lives easier, we can alternatively write the equivalent:

```
1 std::cout << myFavoriteColor->red << "\n";
```

### Maps: std::map

A map is a data structure that maps *keys* to *values*. For example:

```
1 Color TUMBlue = {0, 101, 189};  
2 Color TUMGreen = {162, 173, 0};  
3 Color TUMLightBlue = {100, 160, 200};  
4 std::map<std::string, Color> TUMColorPalette {  
5     {"color-logo", TUMBlue},  
6     {"color-highlight", TUMGreen},  
7     {"color-notes", TUMLightBlue}  
8 }
```

To iterate over a map, we can use a ranged-for loop:

```

1  for( const auto & [name, color] : TUMColorPalette ) {
2      std::cout << "Name: " << name << ", So much red: " << color.red << "\n";
3  }

```

We can also directly access a map element like this:

```

1  TUMColorPalette["color-logo"] = TUMLightBlue;

```

**Important:** If "color-logo" does not already exist in the map, then attempting to access `TUMColorPalette["color-logo"]` (without assigning anything to it) would add an empty value to this key. To check if key exists, use `map.contains(key)` (C++20).

## Advice / C++ Core Guidelines

- R.1: Manage resources automatically using resource handles and RAII
- R.3: A raw pointer (a  $T^*$ ) is non-owning
- R.5: Prefer scoped objects, don't heap-allocate unnecessarily
- R.10: Avoid `malloc()` and `free()`
- R.11: Avoid calling `new` and `delete` explicitly

## Epilogue

In research projects, you will still find raw pointers and explicit calls to `new` and `delete` everywhere. However, actively developed projects typically try to at least implement RAII. The easiest way to do that is using STL containers, such as `std::vector`. Note that raw pointers still have their place in your codebase, and they are often used to pass data to libraries with C-compatible API. By the way, code sanitizers are very useful, but not that widely known and used. Let's change that!